

# Real-Time Collaborative Workspace Backend

Backend Developer Assessment

**Name:** Aaditya Aaryan

**Email:** aadityaaryan639@gmail.com

**Phone:** +91 8340118693

**Date:** December 27, 2025

# Table of Contents

1. Project Overview
2. Architecture Overview
3. Setup & Run Instructions
4. API Documentation
5. Design Decisions & Trade-offs
6. Scalability Considerations
7. Testing Instructions
8. Deployment Instructions
9. Links Summary

# 1. Project Overview

This project is a Real-Time Collaborative Workspace Backend - a production-ready API service that enables teams to collaborate on projects in real-time. It features secure authentication, project and workspace management, role-based access control, and asynchronous code execution jobs.

## Tech Stack:

- \*\*Framework\*\*: FastAPI (Python 3.11)
- \*\*Databases\*\*: PostgreSQL (relational data), MongoDB (job results & logs), Redis (caching & pub/sub)
- \*\*Authentication\*\*: JWT with Argon2 password hashing
- \*\*Async Workers\*\*: Celery for background job processing
- \*\*Real-Time\*\*: WebSocket with Redis Pub/Sub

## Key Features:

- User registration, login, and profile management
- Project CRUD with collaborator invitations
- Workspace management within projects
- Role-based access control (Owner, Collaborator, Viewer)
- Async code execution with status tracking
- Real-time collaboration via WebSockets
- Rate limiting and API caching
- Feature flags for runtime configuration

## 2. Architecture Overview

The system follows a microservices-inspired architecture with clear separation of concerns:

### **System Components:**

1. **FastAPI Application**: Handles HTTP/WebSocket requests, input validation, and routing
2. **PostgreSQL Database**: Stores users, projects, workspaces, and collaborator relationships
3. **MongoDB Database**: Stores high-velocity data like job results and activity logs
4. **Redis**: Serves as cache, rate limiter, pub/sub broker, and Celery message queue
5. **Celery Workers**: Process async jobs like code execution in the background

### **Request Flow:**

1. Client sends request → Load Balancer → FastAPI instance
2. FastAPI validates input, checks auth (JWT), and rate limits (Redis)
3. Business logic executed, data fetched/stored in PostgreSQL/MongoDB
4. For async jobs: Task queued to Redis → Celery worker processes → Result stored in MongoDB
5. For real-time: WebSocket connections use Redis Pub/Sub for cross-instance messaging

### **Database Schema (PostgreSQL):**

- `cw\_users`: User accounts with hashed passwords
- `cw\_projects`: Projects owned by users
- `cw\_workspaces`: Workspaces within projects
- `cw\_collaborators`: Many-to-many relationship with roles

### **Key Libraries:**

- SQLAlchemy 2.0 with async support for PostgreSQL
- Motor for async MongoDB operations
- redis-py for async Redis operations
- passlib + argon2-cffi for password hashing
- PyJWT for token management

## 3. Setup & Run Instructions

### Prerequisites:

- Python 3.11+
- PostgreSQL, MongoDB, Redis (or use Docker)

### Local Setup:

1. Clone the repository:

```
git clone https://github.com/Aadik1ng/Collaborative-Workflow.git  
cd Collaborative-Workflow
```

2. Create virtual environment:

```
python -m venv venv  
source venv/bin/activate # Windows: venv\Scripts\activate
```

3. Install dependencies:

```
pip install -r requirements.txt
```

4. Configure environment variables:

```
cp .env.example .env  
# Edit .env with your database credentials
```

5. Start the API server:

```
uvicorn app.main:app --reload
```

6. Start Celery worker (separate terminal):

```
celery -A app.workers.celery_app worker --loglevel=info
```

### Docker Setup:

```
cd docker  
docker-compose up -d
```

### Environment Variables Required:

- POSTGRES\_URL: PostgreSQL connection string (with +asyncpg)
- MONGODB\_URL: MongoDB connection string
- MONGODB\_DATABASE: Database name
- REDIS\_URL: Redis connection string
- SECRET\_KEY: JWT signing secret
- ALGORITHM: JWT algorithm (HS256)
- ACCESS\_TOKEN\_EXPIRE\_MINUTES: Token lifetime
- REFRESH\_TOKEN\_EXPIRE\_DAYS: Refresh token lifetime

## 4. API Documentation

### Authentication Endpoints

#### POST /api/v1/auth/register

Description: Register new user

Request: { "email": "user@example.com", "username": "user", "password": "SecurePass123!", "full\_name": "John Doe" }

Response: { "id": "uuid", "email": "...", "username": "...", "full\_name": "..." }

Status: 201 Created

#### POST /api/v1/auth/login

Description: Login and get tokens

Request: { "email": "user@example.com", "password": "SecurePass123!" }

Response: { "access\_token": "...", "refresh\_token": "...", "token\_type": "bearer" }

Status: 200 OK

#### POST /api/v1/auth/refresh

Description: Refresh access token

Request: { "refresh\_token": "..." }

Response: { "access\_token": "...", "token\_type": "bearer" }

Status: 200 OK

#### POST /api/v1/auth/logout

Description: Logout user

Request: Header: Authorization: Bearer <token>

Response: { "message": "Logged out successfully" }

Status: 200 OK

#### GET /api/v1/auth/me

Description: Get current user

Request: Header: Authorization: Bearer <token>

Response: { "id": "...", "email": "...", "username": "...", "full\_name": "..." }

Status: 200 OK

#### PUT /api/v1/auth/me

Description: Update profile

Request: { "full\_name": "New Name" }

Response: { "id": "...", "email": "...", "full\_name": "New Name" }

Status: 200 OK

## Project Endpoints

### POST /api/v1/projects

Description: Create project

Request: { "name": "My Project", "description": "...", "is\_public": false}

Response: { "id": "uuid", "name": "...", "owner\_id": "...", "created\_at": "..." }

Status: 201 Created

### GET /api/v1/projects

Description: List projects

Request: Query: ?skip=0&limit=10

Response: [ { "id": "...", "name": "...", ... } ]

Status: 200 OK

### GET /api/v1/projects/{id}

Description: Get project

Request: Path: project ID

Response: { "id": "...", "name": "...", "owner": { ... }, "workspaces": [ ... ] }

Status: 200 OK

### PUT /api/v1/projects/{id}

Description: Update project

Request: { "name": "Updated Name" }

Response: { "id": "...", "name": "Updated Name", ... }

Status: 200 OK

### DELETE /api/v1/projects/{id}

Description: Delete project

Request: Path: project ID

Response: { "message": "Project deleted" }

Status: 200 OK

## Workspace Endpoints

## **POST /api/v1/projects/{id}/workspaces**

Description: Create workspace

Request: { "name": "Workspace 1", "description": "..." }

Response: { "id": "uuid", "name": "...", "project\_id": "..." }

Status: 201 Created

## **GET /api/v1/projects/{id}/workspaces**

Description: List workspaces

Request: Path: project ID

Response: [ { "id": "...", "name": "...", ... } ]

Status: 200 OK

## **Collaborator Endpoints**

### **POST /api/v1/projects/{id}/collaborators**

Description: Invite collaborator

Request: { "email": "collab@example.com", "role": "collaborator" }

Response: { "id": "...", "user\_id": "...", "role": "collaborator" }

Status: 201 Created

### **GET /api/v1/projects/{id}/collaborators**

Description: List collaborators

Request: Path: project ID

Response: [ { "user\_id": "...", "email": "...", "role": "..." } ]

Status: 200 OK

## **Job Endpoints**

### **POST /api/v1/jobs**

Description: Submit code execution job

Request: { "language": "python", "code": "print('Hello')", "timeout": 30 }

Response: { "id": "uuid", "status": "pending", "created\_at": "..." }

Status: 202 Accepted

### **GET /api/v1/jobs/{id}**

Description: Get job status

Request: Path: job ID

**Response:** { "id": "...", "status": "completed", "output": "Hello", "execution\_time": 0.5 }

Status: 200 OK

### **GET /api/v1/jobs**

Description: List user jobs

**Request:** Query: ?skip=0&limit=10

**Response:** [ { "id": "...", "status": "...", ... } ]

Status: 200 OK

### **POST /api/v1/jobs/{id}/cancel**

Description: Cancel job

**Request:** Path: job ID

**Response:** { "message": "Job cancelled" }

Status: 200 OK

## 5. Design Decisions & Trade-offs

### 1. Dual Database Strategy (PostgreSQL + MongoDB)

- Rationale: Relational data (users, projects, roles) benefits from ACID transactions. Non-relational data (job results, logs) needs flexible schemas and high write throughput.
- Trade-off: Increased operational complexity.

### 2. Argon2 for Password Hashing

- Rationale: Winner of Password Hashing Competition, resistant to GPU attacks, no 72-byte limit like bcrypt.
- Trade-off: Slightly higher CPU usage per hash.

### 3. Celery for Async Jobs

- Rationale: Decouples long-running tasks from request cycle, improves API responsiveness.
- Trade-off: Adds Redis as required dependency.

### 4. JWT with Refresh Tokens

- Rationale: Stateless access tokens enable horizontal scaling. Refresh tokens allow session invalidation.
- Trade-off: Requires careful token handling on client.

### 5. Table Name Prefixing (cw\_)

- Rationale: Allows coexistence with other apps in shared database.
- Trade-off: Longer table names.

### 6. Direct argon2-cffi Usage

- Rationale: Bypasses passlib's backend detection issues in serverless environments.
- Trade-off: Less abstraction.

## 6. Scalability Considerations

### Horizontal Scaling:

- FastAPI: Deploy multiple instances behind load balancer (stateless design)
- Celery Workers: Add more workers for increased job throughput
- PostgreSQL: Use read replicas, connection pooling (PgBouncer)
- MongoDB: Sharding for write scaling, replica sets for reads
- Redis: Redis Cluster for HA and scaling

### Performance Optimizations:

- Sliding-window rate limiter protects against abuse
- Redis caching for frequently accessed data
- SQLAlchemy async connection pooling
- Idempotent job processing with unique IDs

### Database Indexing:

- Indexed: user email, username, project owner\_id, workspace project\_id
- TTL indexes on activity logs (7-day expiry)

### Security Measures:

- Argon2 password hashing
- JWT with short-lived access tokens
- Input validation with Pydantic
- CORS configuration
- Rate limiting per IP/user

### Future Enhancements:

- Code execution sandboxing (Docker/gVisor)
- OpenTelemetry for distributed tracing
- Kubernetes deployment with Helm charts

## 7. Testing Instructions

### Running Tests:

```
# Run all tests
pytest

# Run integration tests
pytest tests/integration

# Run unit tests
pytest tests/unit

# Run with coverage
pytest --cov=app --cov-report=term-missing

# Run with verbose output
pytest -v
```

### Test Coverage:

- 58 tests covering authentication, projects, workspaces, collaborators, jobs
- ~59% code coverage focusing on critical paths
- Uses SQLite in-memory for fast integration tests

### Test Categories:

- Unit Tests: Password hashing, JWT tokens, permissions
- Integration Tests: Full API request/response cycles with mocked databases

## 8. Deployment Instructions

### Vercel Deployment:

1. Connect your GitHub repository to Vercel
2. Select "Other" as framework preset
3. Set environment variables in Vercel Dashboard:
  - POSTGRES\_URL (with +asyncpg prefix)
  - MONGODB\_URL
  - MONGODB\_DATABASE
  - REDIS\_URL
  - SECRET\_KEY
  - Other JWT/app settings
4. Override Install Command: pip install -r requirements.txt
5. Deploy - Vercel auto-deploys on push to main

### Important Notes:

- Vercel serverless has timeouts; use persistent hosting for WebSockets
- Celery workers must be deployed separately (Railway, Heroku, VPS)
- Use external databases (Railway, Atlas, Upstash)

### Docker Deployment:

```
cd docker  
docker-compose up -d --build
```

### Railway Deployment:

1. Create PostgreSQL, MongoDB, Redis services
2. Deploy API from GitHub
3. Deploy Celery worker as separate service
4. Configure environment variables

## 9. Links Summary

**GitHub Repository:** <https://github.com/Aadik1ng/Collaborative-Workflow>

**Live Vercel Deployment:** <https://collaborative-workflow.vercel.app>

**Walkthrough Video:** <https://your-video-link.com>